

URL Shortener Project Documentation

Overview

This project entails the development and deployment of a scalable and reliable URL shortener system. The architecture leverages modern cloud-native technologies to ensure high availability, low latency, and the ability to handle significant traffic with zero error rates. Below is a detailed explanation of the technologies used and the reasons behind these choices.

Architecture Components

Frontend

- **Technology:** React
- **Deployment:** AWS Amplify
- **Purpose:**
 - Provides a UI interface for shortening URLs, managing them, and viewing analytics.
 - AWS Amplify was chosen for its simplicity in managing frontend hosting and CI/CD capabilities, as well as its ability to scale automatically based on traffic.
 - URL: [Staging Frontend](#)

Backend

- **Technologies:**
 - **AWS API Gateway:** For routing and managing API endpoints.
 - **AWS Lambda:** For handling backend logic through serverless functions.
 - **DynamoDB (NoSQL):** For storing URL data.
- **Reasoning:**
 - **API Gateway and Lambda:**
 - This serverless architecture ensures automatic scaling, reduced maintenance overhead, and cost efficiency by charging only for actual computation time only!.
 - Functions are separated for better maintainability, with each handling a distinct part of the system:
 - **Shorten URL:** Generates a shortened URL from a long one (working explanation below).
 - **Redirect URL:** Redirects users to the original URL using the shortened link (working explanation below).
 - **Analytics:** Tracks and provides data on URL access (working explanation below).
 - **DynamoDB:**
 - Chosen for its scalability, flexibility, and ability to handle high traffic efficiently.
 - NoSQL databases are well-suited for handling unstructured or semi-structured data, and they scale horizontally, making them ideal for this application, where data access patterns are primarily key-value lookups.

Database Schema

- **Table Name:** ShortenedUrls
- **Primary Key:** shortUrl (String)
- **Attributes:**
 - longUrl (String)

- `clickCount` (Number)
- `createdAt` (String)
- `lastAccessedAt` (String)

Functional Requirements Fulfillment

1. Shorten URL

- **Lambda Function:** `ShortenUrlFunction`
 - **Purpose:** Accepts a long URL and returns a shortened URL.
 - **Workflow:**
 - The function receives the long URL from the frontend via API Gateway.
 - It generates a unique identifier (e.g., an 8-character string) (Hashing Process below).
 - The function constructs a shortened URL using this identifier.
 - It stores the mapping (`shortUrl` to `longUrl`) in **DynamoDB**.
 - Returns the shortened URL to the frontend.
 - **DynamoDB Interaction:**
 - Inserts a new item into the `ShortenedUrls` table with attributes:
 - `shortUrl` : The unique identifier.
 - `longUrl` : The original URL.
 - `clickCount` : Initialized to 0.
 - `createdAt` : Timestamp of creation.
 - **Hash Generation Process**
 - Unique Combination:**
 - Each generated hash combines a timestamp (nanosecond precision) with a unique counter value, ensuring the uniqueness of each hash within the same process.
 - Fast and Efficient:**
 - The use of an atomic counter and timestamp makes the hash generation extremely fast, suitable for high-speed applications that require generating **over a million hashes per second**.
 - URL Safety:**
 - The Base64 URL encoder ensures that the resulting hash does not contain any characters that are unsafe for URLs (`+` , `/` , or `=` are no included).
 - Deterministic Uniqueness:**
 - Within the same process, the combination of the high-resolution timestamp and incrementing counter guarantees that no two hashes will be the same.

2. Redirect URL

- **Lambda Function:** `RedirectUrlFunction`
 - **Purpose:** Redirects users to the original URL when the shortened URL is accessed.
 - **Workflow:**
 - The function receives the `shortUrl` as a path parameter from the API Gateway.
 - It queries **DynamoDB** for the item with the corresponding `shortUrl` .
 - If found, it retrieves the `longUrl` .
 - The `clickCount` attribute is incremented by 1.
 - The function returns the `longUrl` , which is then used to redirect the user.
 - **DynamoDB Interaction:**
 - Reads the item associated with the `shortUrl` .

- Updates the `clickCount` to reflect the number of redirections.

3. Analytics

- **Lambda Function:** `AnalyticsFunction`
 - **Purpose:** Provides analytics, such as the number of times a shortened URL has been accessed.
 - **Workflow:**
 - i. The function receives the `shortUrl` as a path parameter from the API Gateway.
 - ii. It queries **DynamoDB** for the item with the corresponding `shortUrl`.
 - iii. Retrieves and returns analytics data, including `clickCount`, `createdAt`, and `longUrl`.
 - **DynamoDB Interaction:**
 - Reads the item associated with the `shortUrl` to fetch analytics data.

4. Frontend

- **Developed in React** and hosted via **AWS Amplify**, providing a user interface for interacting with the backend API.

5. Backend

- **Managed with API Gateway:**
 - Routes incoming requests to the appropriate Lambda functions.
 - Handles three primary routes:
 - **POST** `url/` : Triggers `ShortenUrlFunction` for URL shortening.
 - **GET** `url/{shortUrl}` : Triggers `RedirectUrlFunction` for redirection.
 - **GET** `url/analytics/{shortUrl}` : Triggers `AnalyticsFunction` for retrieving analytics.

Security Measures

- **URL Validation:** The system checks the format of URLs to prevent malformed.
- **Rate Limiting:** Configured at the API Gateway level to prevent abuse by limiting the number of requests from a single client.

Non-Functional Requirements

- **Scalability:** The system is designed to handle up to **60 concurrent requests per second with a zero-error rate**.
- **Reliability:** Leveraging AWS's managed services ensures high uptime and reliability.
- **Performance:** The use of AWS Lambda and DynamoDB ensures low latency for URL shortening and redirection.

Unimplemented Features

- **User Authentication:** Not yet implemented.
- **CI/CD Pipelines:** Currently, manual deployment processes are used.

Deployment Summary

- **Frontend:** Hosted on AWS Amplify for automated scaling and ease of deployment.
- **Backend:** Deployed using AWS Lambda for serverless computing, managed through API Gateway for routing requests.
- **Database:** DynamoDB used for storing URL mappings and analytics data, chosen for its scalability and ease of integration with Lambda functions.

Future Enhancements

- User Authentication.
- CI/CD Pipelines.
- Tackle ColdStart problem for Lambda Functions.
- Adding User Specific Analytics.